

## Report for Lab 1

### Task 1

... @@ -1,6 +1,6 @@	...
1 IDIR=include	1 IDIR=include
2 #Choose compiler	2 #Choose compiler
3 - #CXX=	3 + CXX=g++
4 CXXFLAGS=-I\$(IDIR) -std=c++11	4 CXXFLAGS=-I\$(IDIR) -std=c++11
5	5
6 ODIR=src	6 ODIR=src
@@ -14,24 +14,20 @@ DEPS = \$(patsubst %, \$(IDIR)/%, \$(DEPS))	
14 _OBJ = deep_core.o vector_ops.o	14 _OBJ = deep_core.o vector_ops.o
15 OBJ = \$(patsubst %, \$(ODIR)/%, \$(OBJ))	15 OBJ = \$(patsubst %, \$(ODIR)/%, \$(OBJ))
16	16
17 - deep_core.o: deep_core.cpp \$(DEPS)	17 + \$(ODIR)/%.o: %.cpp \$(DEPS)
18 - g++ -c -o \$@ \$< \$(CXXFLAGS)	18 + \$(CXX) -c -o \$@ \$< \$(CXXFLAGS)
19	
20 - vector_ops.o: vector_ops.cpp \$(DEPS)	
21 - g++ -c -o \$@ \$< \$(CXXFLAGS)	
22 -	
23	19
24 nnetwork_mpi.o:	20 nnetwork_mpi.o:
25 - g++ -c -o \$@ nnetwork.cxx \$(CXXFLAGS) -DUSE_MPI	21 + \$(CXX) -c -o \$@ nnetwork.cxx \$(CXXFLAGS) -DUSE_MPI
26	22
27 nnetwork.o:	23 nnetwork.o:
28 - g++ -c -o \$@ nnetwork.cxx \$(CXXFLAGS)	24 + \$(CXX) -c -o \$@ nnetwork.cxx \$(CXXFLAGS)
29	25
30 nnetwork_mpi: \$(OBJ) nnetwork_mpi.o	26 nnetwork_mpi: \$(OBJ) nnetwork_mpi.o
31 - g++ -o \$@ \$^ \$(LIBS)	27 + \$(CXX) -o \$@ \$^ \$(LIBS)
32	28
33 nnetwork: \$(OBJ) nnetwork.o	29 nnetwork: \$(OBJ) nnetwork.o
34 - g++ -o \$@ \$^ \$(LIBS)	30 + \$(CXX) -o \$@ \$^ \$(LIBS)
35	31
36 run_serial:	32 run_serial:
37 ./nnetwork	33 ./nnetwork

### Task 2

(a) Make debug build:

... @@ -1,7 +1,7 @@	
1 IDIR=include	1 IDIR=include
2 #Choose compiler	2 #Choose compiler
3 CXX=g++	3 CXX=g++
4 - CXXFLAGS=-I\$(IDIR) -std=c++11	4 + CXXFLAGS=-I\$(IDIR) -std=c++11 -g
5	5
6 ODIR=src	6 ODIR=src
7 IDIR=./lib	7 IDIR=./lib

(b) Fix segfault:

@@ -115,7 +115,7 @@ int main(int argc, char * argv[]) {	
115 print ( yhat, 10, 10 );	115 print ( yhat, 10, 10 );
116 cout << "Ground truth:" << "\n";	116 cout << "Ground truth:" << "\n";
117 print ( b_y, 10, 10 );	117 print ( b_y, 10, 10 );
118 - vector<float> loss_m //=- yhat - b_y;	118 + vector<float> loss_m = yhat - b_y;
119 float loss = 0.0;	119 float loss = 0.0;
120 for (unsigned k = 0; k < BATCH_SIZE*10;	120 for (unsigned k = 0; k < BATCH_SIZE*10;
++k){	++k){
loss += loss_m[k]*loss_m[k];	loss += loss_m[k]*loss_m[k];
121	121

(c) After variable "size" is initialized, the value it contains is 32768. z is a const reference vector, containing floats.  $z[3] = 1.96690202$ .

```
Breakpoint 1, relu (z=...) at src/deep_core.cpp:77
77 vector<float> output;
(gdb) l
72 return output;
73 }
74
75 vector<float> relu(const vector<float>& z){
76 int size = z.size();
77 vector<float> output;
78 for( int i = 0; i < size; ++i ) {
79 if (z[i] < 0){
80 output.push_back(0.0);
81 }
(gdb) p size
$1 = 32768
(gdb) p z[3]
$2 = (const __gnu_cxx::__alloc_traits<std::allocator<float>, float>::value_type &) @0x555555f20dc: 1.96690202
```

### Task 3

- (a) First, we have created a rule in the *Makefile* to enable performance profiling. In particular, the rule is given by

```
run_perf:
    perf record ./nnetwork .
```

- (b) After running the added rule, we have inspected the generated performance file `perf.data` by running the command `perf report`. Overall, there are three time-consuming operations.

- (1) The most time is spent in the `dot` function in the file `vector_ops.cpp`. It requires about 67.62% of running time and computes the product of two matrices. Optimizing it could highly improve the total running time of the training.
- (2) The second and third most time-consuming operations are array accesses on vector data-structures with 18.41% and 9.36% respectively. These operations might be optimized by reducing the number of cache misses while accessing vector elements.
- (3) Apart from those three operations, there are not any others with significant share on the running time.

- (c) To monitor the number of LLC cache misses, we have run the command `perf stat -e LLC-load-misses ./nnetwork`. However, the machine that we have been using did not support for that. The program returned

Performance counter stats for './nnetwork':

```
<not supported>      LLC-load-misses:u
```

```
412.636442944 seconds time elapsed
```

```
397.622062000 seconds user
```

```
6.620183000 seconds sys .
```