

## Report for Lab 1

### Task 1

The attached screenshot shows the changes made to introduce a variable for the C++ compiler used.

1	IDIR=include	1	IDIR=include
2	#Choose compiler	2	#Choose compiler
3	- #CXX=	3	+ CXX=g++
4	CXXFLAGS=-I\$(IDIR) -std=c++11	4	CXXFLAGS=-I\$(IDIR) -std=c++11
5		5	
6	ODIR=src	6	ODIR=src
14	_OBJ = deep_core.o vector_ops.o	14	_OBJ = deep_core.o vector_ops.o
15	OBJ = \$(patsubst %, \$(ODIR)/%, \$(OBJ))	15	OBJ = \$(patsubst %, \$(ODIR)/%, \$(OBJ))
16		16	
17	- deep_core.o: deep_core.cpp \$(DEPS)	17	+ \$(ODIR)/%.o: %.cpp \$(DEPS)
18	- g++ -c -o \$@ \$< \$(CXXFLAGS)	18	+ \$(CXX) -c -o \$@ \$< \$(CXXFLAGS)
19	-		
20	- vector_ops.o: vector_ops.cpp \$(DEPS)		
21	- g++ -c -o \$@ \$< \$(CXXFLAGS)		
22	-		
23		19	nnetwork_mpi.o:
24	nnetwork_mpi.o:	20	+ \$(CXX) -c -o \$@ nnetwork.cxx \$(CXXFLAGS) -DUSE_MPI
25	- g++ -c -o \$@ nnetwork.cxx \$(CXXFLAGS) -DUSE_MPI	21	
26		22	nnetwork.o:
27	nnetwork.o:	23	+ \$(CXX) -c -o \$@ nnetwork.cxx \$(CXXFLAGS)
28	- g++ -c -o \$@ nnetwork.cxx \$(CXXFLAGS)	24	
29		25	nnetwork_mpi: \$(OBJ) nnetwork_mpi.o
30	nnetwork_mpi: \$(OBJ) nnetwork_mpi.o	26	+ \$(CXX) -o \$@ \$^ \$(LIBS)
31	- g++ -o \$@ \$^ \$(LIBS)	27	
32		28	nnetwork: \$(OBJ) nnetwork.o
33	nnetwork: \$(OBJ) nnetwork.o	29	+ \$(CXX) -o \$@ \$^ \$(LIBS)
34	- g++ -o \$@ \$^ \$(LIBS)	30	
35		31	run_serial:
36	run_serial:	32	./nnetwork
37	./nnetwork	33	

### Task 2

(a) To make a debug build, we added the following flag.

1	IDIR=include	1	IDIR=include
2	#Choose compiler	2	#Choose compiler
3	CXX=g++	3	CXX=g++
4	- CXXFLAGS=-I\$(IDIR) -std=c++11	4	+ CXXFLAGS=-I\$(IDIR) -std=c++11 -g
5		5	
6	ODIR=src	6	ODIR=src
7	IDIR = ../Lib	7	IDIR = ../Lib

(b) We changed the following lines to fix the *segmentation* fault.

115	print ( yhat, 10, 10 );	115	print ( yhat, 10, 10 );
116	cout << "Ground truth:" << "\n";	116	cout << "Ground truth:" << "\n";
117	print ( b_y, 10, 10 );	117	print ( b_y, 10, 10 );
118	- vector<float> loss_m //=- yhat - b_y;	118	+ vector<float> loss_m = yhat - b_y;
119	float loss = 0.0;	119	float loss = 0.0;
120	for (unsigned k = 0; k < BATCH_SIZE*10;	120	for (unsigned k = 0; k < BATCH_SIZE*10;
121	++k){	121	++k){
	loss += loss_m[k]*loss_m[k];		loss += loss_m[k]*loss_m[k];

(c) After variable *size* is initialized, the value it contains is 32768. The variable *z* is a *const* reference vector, containing floats.  $z[3] = 1.96690202$ .

```

Breakpoint 1, relu (z=...) at src/deep_core.cpp:77
77     vector<float> output;
(gdb) l
72         return output;
73     }
74
75     vector<float> relu(const vector<float>& z){
76         int size = z.size();
77         vector<float> output;
78         for( int i = 0; i < size; ++i ) {
79             if (z[i] < 0){
80                 output.push_back(0.0);
81             }
(gdb) p size
$1 = 32768
(gdb) p z[3]
$2 = (const __gnu_cxx::__alloc_traits<std::allocator<float>, float>::value_type &) @0x555555f20dc: 1.96690202

```

### Task 3

- (a) First, we have created a rule in the *Makefile* to enable performance profiling. In particular, the rule is given by

```

run_perf:
    perf record ./nnetwork .

```

- (b) After running the added rule, we have inspected the generated performance file `perf.data` by running the command `perf report`. Overall, there are three time-consuming operations.

- (1) The most time is spent in the `dot` function in the file `vector_ops.cpp`. It requires about 67.62% of running time and computes the product of two matrices. Optimizing it could highly improve the total running time of the training.
- (2) The second and third most time-consuming operations are array accesses on vector data-structures with 18.41% and 9.36% respectively. These operations might be optimized by reducing the number of cache misses while accessing vector elements.
- (3) Apart from those three operations, there are not any others with significant share on the running time.

- (c) To monitor the number of LLC cache misses, we have run the command `perf stat -e LLC-load-misses ./nnetwork`. The program returned

Performance counter stats for './nnetwork':

2,043,328 LLC-load-misses:u

215.866059446 seconds time elapsed

215.823836000 seconds user

0.035999000 seconds sys