

## Report for Lab 4

### Task 1 – OpenMP parallel for

```
#pragma omp parallel
{
    int thread_id = omp_get_thread_num();
    int tcount = omp_get_num_threads();

    std::cout << "Thread ID " << thread_id << std::endl;
    if (thread_id == 0) std::cout << "Total threads: " << tcount << std::endl;

    #pragma omp for
    for ( ... ) {
        ...
    }
}
```

1. The thread id inside the parallel region is printed out with the aid of the `omp_get_thread_num()` function.
2. The total number of threads inside the parallel region is printed out with the aid of the `omp_get_num_threads()` function. It is printed only if `thread_id == 0`. Therefore, it is output only once.

### Task 2 – Parallel performance gains

1. Text...
2. Text...

### Task 3 – OpenMP scheduling policies

	(no param)	1	16	128
<code>schedule(static, &lt;param&gt;)</code>	4.6989 s	7.6719 s	4.7186 s	6.2623 s
<code>schedule(dynamic, &lt;param&gt;)</code>	10.9646 s	10.8904 s	4.8008 s	6.3657 s
<code>schedule(guided, &lt;param&gt;)</code>	4.8178 s	4.8473 s	4.7503 s	6.3849 s

**Table 1:** Total running times for all loop scheduling strategies.

1. We have included all three scheduling strategies, namely `static`, `dynamic`, and `guided`, in our experiments.
2. Table 1 shows the running times of all three strategies with either no parameter specified, a chunk size of 1, 16, or 128.

3. The optimal policy is `schedule(static)`. On the one hand, since all loop iterations of the GEMM loops take more or less the same time, there is no benefit of using `dynamic` scheduling because the extra time spend for scheduling does not improve the distribution of work significantly. On the other hand, `schedule(static)` is more cache-locality preserving than the `guided` scheduling strategy wherefore the `static` policy wins the comparison.

#### Task 4 – Which loop to parallelize?

1. Text...
2. Text...
3. Text...

#### Task 5 – Exploring SIMD

1. Without any loop permutations, there is not any vectorizing happening for any of the loops in the GEMM function.
2. With the loop permutation from Lab 2, the optimization flags `-O1` and `-O2` yield also no vectorization of the loops. However, using the optimization flag `-O3` vectorizes the inner-most loop (column loop; `src/vector_ops.cpp:242:39: note: loop vectorized`). The second inner-most loop ( $k$  loop), however, is not vectorized (`src/vector_ops.cpp:241:31: note: bad data references.`). Also, the outer-most loop (row loop) is not vectorized either (`src/vector_ops.cpp:240:31: note: not vectorized: multiple nested loops.`).

3. With no vectorized GEMM loops (optimization flag `-O2`), the program has the following running time.

```
Total time in dot: 30.82s
Total time (all): 33.5521s
Percentage in dot: 91.8574%
```

With the vectorized inner-most loop (optimization flag `-O3`), the program's running time improves as seen below.

```
Total time in dot: 10.9205s
Total time (all): 13.4652s
Percentage in dot: 81.1018%
```

The total speedup is  $S = \frac{33.5521\text{s}}{13.4652\text{s}} = 2.4918$ .