

Report for Lab 4

Task 1 – OpenMP parallel for

```
#pragma omp parallel
{
    int thread_id = omp_get_thread_num();
    int tcount = omp_get_num_threads();

    std::cout << "Thread ID " << thread_id << std::endl;
    if (thread_id == 0) std::cout << "Total threads: " << tcount << std::endl;

    #pragma omp for
    for ( ... ) {
        ...
    }
}
```

1. The thread id inside the parallel region is printed out with the aid of the `omp_get_thread_num()` function.
2. The total number of threads inside the parallel region is printed out with the aid of the `omp_get_num_threads()` function. It is printed only if `thread_id == 0`. Therefore, it is output only once.

Task 2 – Parallel performance gains

1. We use the `time` command to time the total execution time of the program and use `omp_get_wtime()` to time the parallel execution.

```
double total_time_in_parallel = 0; // variable in global namespace
double start = omp_get_wtime();
... parallel code ...
double end = omp_get_wtime();
total_time_in_parallel += end - start;
```

The parallel execution time decreases when increasing number of threads, the table below shows the results.

2. The table below shows the execution time and efficiency:

	Serial	1 Thread	2 Threads	4 Threads
Total time	13.4743 s	10.5384 s	6.6796 s	4.6443 s
Efficiency	0.7821	1.0000	0.7889	0.5673

Table 1: Efficiency for different number of threads

Task 3 – OpenMP scheduling policies

	(no param)	1	16	128
<code>schedule(static, <param>)</code>	4.6989 s	7.6719 s	4.7186 s	6.2623 s
<code>schedule(dynamic, <param>)</code>	10.9646 s	10.8904 s	4.8008 s	6.3657 s
<code>schedule(guided, <param>)</code>	4.8178 s	4.8473 s	4.7503 s	6.3849 s

Table 2: Total running times for all loop scheduling strategies.

1. We have included all three scheduling strategies, namely **static**, **dynamic**, and **guided**, in our experiments.
2. Table 2 shows the running times of all three strategies with either no parameter specified, a chunk size of 1, 16, or 128.
3. The optimal policy is **schedule(static)**. On the one hand, since all loop iterations of the GEMM loops take more or less the same time, there is no benefit of using **dynamic** scheduling because the extra time spend for scheduling does not improve the distribution of work significantly. On the other hand, **schedule(static)** is more cache-locality preserving than the **guided** scheduling strategy wherefore the **static** policy wins the comparison.

Task 4 – Which loop to parallelize?

All of task 4 is done using 4 threads.

1. j loop(column) parallelism implementation:

```
for( int row = 0; row < m1_rows; ++row ) {
    #pragma omp parallel for shared(row) //task 4.1
    for( int col = 0; col < m2_columns; ++col ) {
        for( int k = 0; k < m1_columns; ++k ) {
            output[ row * m2_columns + col ] += ... ;
        }
    }
}
```

2. For k loop(entry) parallelism implementation, each output vector element will be written simultaneously by all threads, therefore we decided to use reduction on a local variable and copy the value to each vector element when parallel execution is done:

```
for( int row = 0; row < m1_rows; ++row ) {
    for( int col = 0; col < m2_columns; ++col ) {
        float sum = 0;
        #pragma omp parallel for shared(row, col) reduction(+: sum)
        for( int k = 0; k < m1_columns; ++k ) {
            sum += m1[ row * m1_columns + k ] * m2[ k * m2_columns + col ];
        }
        output[ row * m2_columns + col ] = sum;
    }
}
```

3. Table below shows the parallel execution time of the 3 strategies.

Apparently, parallelizing the k loop is the worst strategy because the critical region basically eliminates the benefits of multithreading, and even worse, the thread synchronization added extra cost.

Parallelizing the i loop(row) slightly outperforms parallelizing j loop(column), which can be explained by cache false sharing. Compared with the i loop parallelism, the j loop parallelism means output vector accesses by each thread is more interleaved, i.e. higher chance of false sharing, and also the j loop introduce more thread management overhead because the parallel section is called `m1_rows` times per function call.

Parallelized loop:	i	j	k
Parallel execution time:	17.8018 s	19.7789 s	186.6380 s

Table 3: performance for different strategies

Task 5 – Exploring SIMD

- Without any loop permutations, there is not any vectorizing happening for any of the loops in the GEMM function.
- With the loop permutation from Lab 2, the optimization flags `-O1` and `-O2` yield also no vectorization of the loops. However, using the optimization flag `-O3` vectorizes the inner-most loop (column loop; `src/vector_ops.cpp:242:39: note: loop vectorized`). The second inner-most loop (k loop), however, is not vectorized (`src/vector_ops.cpp:241:31: note: bad data references.`). Also, the outer-most loop (row loop) is not vectorized either (`src/vector_ops.cpp:240:31: note: not vectorized: multiple nested loops.`).
- With no vectorized GEMM loops (optimization flag `-O2`), the program has the following running time.

```
Total time in dot: 30.82s
Total time (all): 33.5521s
Percentage in dot: 91.8574%
```

With the vectorized inner-most loop (optimization flag `-O3`), the program's running time improves as seen below.

```
Total time in dot: 10.9205s
Total time (all): 13.4652s
Percentage in dot: 81.1018%
```

The total speedup is $S = \frac{33.5521\text{s}}{13.4652\text{s}} = 2.4918$.