

## Report for Lab 6

### Task 1 – Basic Implementation

- The code below shows our implementation:

```

a = (float*)malloc(sizeof(float)* n * n);
b = (float*)malloc(sizeof(float)* n * n);
c = (float*)malloc(sizeof(float)* n * n);
d = (float*)malloc(sizeof(float)* n * n);
...
/* Task: Memory Allocation */
cudaMalloc((void **) &cuda_a, size);
cudaMalloc((void **) &cuda_b, size);
cudaMalloc((void **) &cuda_c, size);
cudaMemset(&cuda_c, 0, size);

/* Task: CUDA Memory Copy from Host to Device */
cudaMemcpy(cuda_a, a, size, cudaMemcpyHostToDevice);
cudaMemcpy(cuda_b, b, size, cudaMemcpyHostToDevice);

/* Task: Number of Blocks and Threads && Dimention*/
dim3 dimGrid(100, 100, 1);
dim3 dimBlock(10, 10, 1);

// Kernel Execution
matMultCUDA << < dimGrid, dimBlock >> >(cuda_a , cuda_b , cuda_c , n);

/* Task: CUDA Memory Copy from Device to Host */
cudaMemcpy(c, cuda_c, size, cudaMemcpyDeviceToHost);

```

- For each thread we calculate the global unique thread id and use it to determine which element in the output matrix the thread shall calculate. Kernel implementation is as follow:

```

__global__ static void matMultCUDA(const float* a, const float* b, float* c, int n)
{
    int threadId = (threadIdx.y + blockIdx.y * blockDim.y) * gridDim.x * blockDim.x +
                  (threadIdx.x + blockIdx.x * blockDim.x);
    int row = threadId/n;
    int col = threadId%n;
    int elementIdx = row * n + col;
    if(elementIdx < n*n) {
        for (int i = 0; i < n; i++) {
            c[elementIdx] += a[row * n + i] * b[i * n + col];
        }
    }
}

```

- We tried the following configurations(all Z-dimension settings are 1) in Table 1 :

	Grid_X	Grid_Y	Block_X	Block_Y
Config 1	1000	1000	1	1
Config 2	125	125	8	8
Config 3	50	50	20	20
Config 4	50	25	40	20

**Table 1:** 4 configurations

- we used `nvcc -o matmul_v1 cuda_mm.cu; ./matmul_v1` to compile and run the program.
- The results is as shown in Table 2. Since warp size is 32, thread block with fewer than 32 threads will not fully utilize the hardware. From the results, it is obvious that thread block with only a single thread is considerably worse than other configurations. Other configurations show similar performance in terms of execution time.

	Config 1	Config 2	Config 3	Config 4
Execution Time	3.860 s	3.168 s	3.183 s	3.144 s

**Table 2:** Execution Time

## Task 2 – Tiling Implementation using Shared Memory

- The used machine has 49152 bytes shared memory per thread block.
- Since we are using `floats` as a data type, each thread block can store 12288 `floats` in its shared memory. Therefore, each of the shared-memory matrices (`a`, `b` and `c`) can contain up to 4096 entries.
- The following snippet shows the implementation of the new tiling matrix multiplication kernel.

```
__global__ static void matMultCUDA(const float *a, const float *b, float *c, int n)
{
    __shared__ float shared_a_tile[BLOCK_SIZE * BLOCK_SIZE];
    __shared__ float shared_b_tile[BLOCK_SIZE * BLOCK_SIZE];
    __shared__ float shared_c_tile[BLOCK_SIZE * BLOCK_SIZE];

    int j = threadIdx.x + blockIdx.x * BLOCK_SIZE;
    int i = threadIdx.y + blockIdx.y * BLOCK_SIZE;

    if (0 <= i && i < n && 0 <= j && j < n)
    {
        // Init shared output tile
```

```

shared_c_tile[threadIdx.y * BLOCK_SIZE + threadIdx.x] = 0;

// Iterate over all tiles
for (int k_tile = 0; k_tile < n; k_tile += BLOCK_SIZE)
{
    // Copy current tiles
    shared_a_tile[threadIdx.y * BLOCK_SIZE + threadIdx.x] =
        a[i * n + (threadIdx.x + k_tile)];
    shared_b_tile[threadIdx.y * BLOCK_SIZE + threadIdx.x] =
        b[(threadIdx.y + k_tile) * n + j];

    __syncthreads();

    // Compute tile of elements
    for (int k = k_tile; k < k_tile + BLOCK_SIZE && k < n; ++k)
    {
        shared_c_tile[threadIdx.y * BLOCK_SIZE + threadIdx.x] +=
            shared_a_tile[threadIdx.y * BLOCK_SIZE + (k % BLOCK_SIZE)]
            * shared_b_tile[(k % BLOCK_SIZE) * BLOCK_SIZE + threadIdx.x];
    }

    __syncthreads();
}

c[i * n + j] = shared_c_tile[threadIdx.y * BLOCK_SIZE + threadIdx.x];
}
}

```

- The BLOCK\_SIZE variable in the following snippet is instantiated with either 4, 8, 16 or 32. It is limited to 32 since  $32 \cdot 32 = 1024$  is the maximum of threads per thread block. In the case of a block size of  $4 \times 4$ , for example, there are  $\text{ceil}(n / \text{BLOCK\_SIZE}) = 250$  blocks per dimension ( $250 \times 250$  blocks in total).

```

// 2D config with given block size
dim3 dimGrid(ceil(n / BLOCK_SIZE), ceil(n / BLOCK_SIZE), 1);
dim3 dimBlock(BLOCK_SIZE, BLOCK_SIZE, 1);

```

- Table 3 shows the execution times from 5 runs for all different configurations. Since the differences are not significant, there is no configuration that is better suited than the others. However, we expected that blocks of size  $32 \times 32$  are best suited for the tiling matrix kernel since the larger the block size is, the less `__syncthreads()` calls there are, which means less overhead.

	4	8	16	32
Run 1	3.252 s	3.189 s	3.323 s	3.085 s
Run 2	3.138 s	3.183 s	3.285 s	3.355 s
Run 3	3.127 s	3.115 s	3.111 s	3.317 s
Run 4	3.123 s	3.288 s	3.097 s	3.100 s
Run 5	3.140 s	3.139 s	3.094 s	3.104 s
Average	3.156 s	3.183 s	3.182 s	3.192 s

**Table 3:** Execution times for different configurations.